Universal Domain Adaptation through Self Supervision (GNR 650)

Uddeshya Singh (22m2152) Sachin Giroh(22m2159)

Motivation

- Deep neural networks can learn highly discriminative representations for image recognition tasks, but they do not generalize well to new domains that are not distributed identically to the training data.
- Domain adaptation (DA) aims to transfer representations of source categories to novel target domains without additional supervision.
- Traditional unsupervised domain adaptation methods assume that all source categories are present in the target domain. However, in practice, little may be known about the category overlap between the two domains.

Novelties

- The paper introduces a new framework called "Domain Adaptative Neighborhood Clustering via Entropy optimization (DANCE)".
- DANCE combines two novel ideas:
    1. A neighborhood clustering technique to learn the structure of the target domain in a self-supervised way since source categories cannot be fully relied upon.
    2. Entropy-based feature alignment and rejection. This aligns target features with the source or rejects them as unknown categories based on their entropy.

Major Contributions

The paper introduces an innovative approach to universal domain adaptation, adeptly tackling various DA scenarios (open, close, partial, open/partial) without prior knowledge of category distribution. Its strength lies in the unique idea of clustering each target sample to a source class prototype or its target domain neighbor. Comprehensive empirical evaluations across diverse datasets and domain shifts underscore its superiority, outperforming existing methods significantly in close-set DA, partial DA, and open-set DA.

Critical Analysis

Strengths

- The paper introduces a novel approach to universal domain adaptation through the design of neighborhood clustering and entropy separation loss. This fresh attempt distinctly sets it apart.
- DANCE consistently outperforms competing methods, notably in the realms of CDA, ODA, PDA, and OPDA.
- The research emphasizes a comprehensive adaptation setting, presenting a singular solution fortified by two specifically crafted losses.
- The claims are robustly backed by a wealth of experimental evidence.

Limitations

- The paper's articulation, especially in Section 3 and its figures, could benefit from enhanced clarity.
- The densely packed experimental section can be challenging to navigate given its breadth relative to page constraints
- While the methods pivot primarily on empirical observations, the DA problem's inherent complexity as a machine learning issue warrants theoretical justifications or explanations which are currently absent.